



Cassandra

Alpesh Vasant, G Academy


Introduction To Cassandra

- It's a Wide Column NoSQL database
- Open Source- So no paid support
- Distributed – Data stored in clusters, data centers and on nodes
- Handles large amount of data may be spreaded physically across world
- Hence provides services with very high available period

Introduction To Cassandra


- **Apple** uses 100,000 Cassandra nodes
- **Netflix** uses Cassandra as their back-end database for their streaming services
- **SoundCloud** uses Cassandra to store the dashboard of their users
- **Uber** uses Cassandra to store around 10000 features in their daily updated company-wide Features
- **Cisco's Webex** uses Cassandra to store user feed and activity in near real time
- **Discord** switched to Cassandra to store billions of messages from MongoDB in November 2005

Architecture Of Cassandra




- Token ring like architecture with no master no slave
- Nodes are logically distributed like ring. Hash values of the keys are used to distribute the data among nodes in the cluster.
- Automatic data distribution across all nodes

Architecture Of Cassandra



- Node: Base component. Data is stored on Node
- Data Centre: Collection of Nodes is called Data Centers
- Cluster: Collections of many data Centers
- Architecture supports multiple data centers
- Data can be replicated across data centers

Architecture Of Cassandra



- Commit Log: Every WRITE operation is written here. Used for crash recovery
- Mem Table: Post data written in commit log, data is also written in Mem table but temporally
- SS Table: Once the Mem Table reaches to certain preset threshold, data is shifted to SS Table (disk file)

Comparison between Cassandra and MongoDB

MongoDB	Cassandra
Document based NoSQL database	Wide Column NoSQL database
Data is stored in documents in the form of Key-Value pair	Stores data in the form of traditional rows and columns
Explicit data type is not required to declare while storing yet it stores large type of data including: strings, numbers, geo data, dates, arrays, decimal, nested objects, and binary data	Each column has a specific data type that it can store that has to be specified at the time of table creation
Compare to Cassandra availability is bit compromised due to master slave model	Cassandra provides very high availability due to token ring type of architecture
JSON is the language for query	CQL is language for query
Has built in aggregation framework to run ETL pipeline	No built in aggregation framework. If needed, it used external tools like Apache Hadoop or Apache Spark

Features Of Cassandra

- **Elastic scalability** – Cassandra is highly scalable; it allows to add more hardware to accommodate more customers and more data as per requirement.
- **Always on architecture** – Cassandra is continuously available for business-critical applications that cannot afford a failure.
- **Fast linear-scale performance** – As you increase the number of nodes in the cluster, you handle more simultaneous requests by customers. Therefore it maintains a quick response time.
- **Flexible data storage** – Cassandra accommodates all possible data formats including: structured, semi-structured, and unstructured. It can dynamically accommodate changes to your data structures according to your need.

Features Of Cassandra

- **Easy data distribution** – Cassandra provides the flexibility to distribute data where you need by replicating data across multiple data centers.
- **Transaction support** – Cassandra supports properties like Atomicity, Consistency, Isolation, and Durability (ACID).
- **Fast writes** – Cassandra was designed to run on cheap commodity hardware. It performs blazingly fast writes and can store hundreds of terabytes of data, without sacrificing the read efficiency.